

LECTURE 7 LARGE SAMPLE THEORY

Limits and convergence concepts: almost sure, in probability and in mean

Let $\{a_n : n = 1, 2, \dots\}$ be a sequence of non-random real numbers. We say that a is the limit of $\{a_n\}$ if for *all* real $\delta > 0$ we can find an integer N_δ such that for *all* $n \geq N_\delta$ we have that $|a_n - a| < \delta$. When the limit exists, we say that $\{a_n\}$ converges to a , and write $a_n \rightarrow a$ or $\lim_{n \rightarrow \infty} a_n = a$. In this case, we can make the elements of $\{a_n : n \geq N\}$ arbitrarily close to a by choosing N sufficiently large. Naturally, if $a_n \rightarrow a$, we have that $a_n - a \rightarrow 0$.

The concept can be extended to vectors or matrices as well. Let $\{A_n : n = 1, 2, \dots\}$ be a $m \times k$ matrix. Then $A_n \rightarrow A$ if for all $i = 1, \dots, m$ and $j = 1, \dots, k$ we have that the (i, j) -th element of A_n converges to the (i, j) -th element of A .

The concept of convergence cannot be applied in a straightforward way to sequences of random variables. This is so because a random variable is a *function* from the sample space Ω to the real line. The solution is to consider convergence of a *non-random* sequence derived from the random one. Since there are many ways to derive non-random sequences, there exist several stochastic convergence concepts. Let $\{X_n : n = 1, 2, \dots\}$ be a sequence of *random* variables. Let X be random or non-random (i.e. it is possible that $X(\omega)$ is the same for all $\omega \in \Omega$). We will consider *non-random* sequences with the following typical elements: (i) $X_n(\omega) - X(\omega)$ for a fixed $\omega \in \Omega$, (ii) $E|X_n - X|^r$, and (iii) $\Pr(|X_n - X| > \varepsilon)$ for some $\varepsilon > 0$. These are sequences of non-random real numbers, and, consequently, the usual definition of convergence applies to each of them leading to a corresponding definition of stochastic convergence:

- (i) **Almost sure (with probability one or pointwise) convergence.** We say that X_n converges to X almost surely if $\Pr(\{\omega : X_n(\omega) \rightarrow X(\omega)\}) = 1$ as $n \rightarrow \infty$.
- (ii) **Convergence in r -th mean.** X_n converges to X in r -th mean if $E|X_n - X|^r \rightarrow 0$ as $n \rightarrow \infty$.
- (iii) **Convergence in probability.** X_n converges in probability to X if for all $\varepsilon > 0$, $\Pr(|X_n - X| \geq \varepsilon) \rightarrow 0$ as $n \rightarrow \infty$. It is denoted as $X_n \rightarrow_p X$ or $\text{plim } X_n = X$. Alternatively, convergence in probability can be defined as $\Pr(|X_n - X| < \varepsilon) \rightarrow 1$ for all $\varepsilon > 0$. The two definitions are equivalent.

The almost sure convergence is the convergence concept most closely related to that of non-random sequences. It implies that for almost all outcomes of the random experiment, X_n converges to X . Convergence in probability is the most useful among the three concepts, since many of the results in econometrics use this type of convergence.

One can show that $X_n \rightarrow X$ almost surely if and only if for all $\varepsilon > 0$,

$$\lim_{n \rightarrow \infty} \Pr\left(\sup_{m \geq n} |X_m - X| \geq \varepsilon\right) = 0. \quad (1)$$

Thus, while convergence in probability focuses only on the marginal distribution of $|X_n - X|$ as $n \rightarrow \infty$, almost sure convergence puts restriction on the joint behavior of all random elements in the sequence $|X_n - X|, |X_{n+1} - X|, \dots$

Almost sure convergence and convergence in r -th mean each imply convergence in probability. However, converse does not hold.

Example. (a) Consider a random variable X_n such that $\Pr(X_n = 0) = 1 - 1/n$, and $\Pr(X_n = n) = 1/n$.

In this example, $X_n \rightarrow_p 0$, since for $\varepsilon > 0$,¹

$$\begin{aligned} \Pr(|X_n| \geq \varepsilon) &\leq \Pr(X_n = n) \\ &= 1/n \\ &\rightarrow 0. \end{aligned} \tag{2}$$

However,

$$E|X_n| = 1,$$

and, consequently, X_n does not converge in mean to zero.

(b) In part (a) of the example, only the marginal distributions of the elements of $\{X_n\}$ have been described. In order to discuss almost sure convergence, we need to characterize the joint distribution of the elements of the sequence $\{X_n\}$. One simple model for X_n in part (a) is to let ω be drawn from the uniform distribution on $[0, 1]$, so that $\Pr(a \leq \omega \leq b) = |b - a|$ for $a, b \in [0, 1]$. Now,

$$X_n(\omega) = \begin{cases} n, & \omega \in [0, 1/n], \\ 0, & \omega \in [1/n, 1], \end{cases}$$

so that $\Pr(X_n = 0) = 1 - 1/n$ and $\Pr(X_n = n) = 1/n$ as before. Using (1) one can show that in this example $X_n \rightarrow X$ almost surely. To show this, first note that the condition (1) can be stated equivalently as $\lim_{n \rightarrow \infty} \Pr(\sup_{m \geq n} |X_m - X| < \varepsilon) = 1$ for all $\varepsilon > 0$. Now, with the limit $X = 0$, define

$$P_{n,\varepsilon} = \Pr\left(\sup_{m \geq n} |X_m| < \varepsilon\right).$$

We have that for all $\varepsilon > 0$ that are small enough,

$$\begin{aligned} P_{n,\varepsilon} &= \Pr(|X_n| < \varepsilon, |X_{n+1}| < \varepsilon, \dots) \\ &= \Pr(X_n = 0, X_{n+1} = 0, \dots) \\ &= \Pr\left(\omega \geq \frac{1}{n}, \omega \geq \frac{1}{n+1}, \dots\right) \\ &= \Pr\left(\omega \geq \frac{1}{n}\right) \\ &= 1 - \frac{1}{n}. \end{aligned}$$

Thus, for all $\varepsilon > 0$ and as $n \rightarrow \infty$,

$$\begin{aligned} \lim_{n \rightarrow \infty} \Pr\left(\sup_{m \geq n} |X_m - X| < \varepsilon\right) &= \lim_{n \rightarrow \infty} P_{n,\varepsilon} \\ &= \lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right) \\ &= 1. \end{aligned}$$

(c) Now, suppose instead that the elements of the sequence $\{X_n\}$ are independent, but continue to

¹Note that if one chooses ε such that $\varepsilon > n$, $\Pr(|X_n| \geq \varepsilon)$ is obviously zero in this example. That is why we have “ $\leq$ ” in (2). Note also that, when establishing convergence results of this kind, we are concerned with small ε ’s.

assume that $\Pr(X_n = 0) = 1 - 1/n$ and $\Pr(X_n = n) = 1/n$. For all positive integers n , we have now that

$$\begin{aligned}
P_{n,\varepsilon} &= \Pr(X_n = 0, X_{n+1} = 0, \dots) \\
&= \prod_{m=n}^{\infty} \left(1 - \frac{1}{m}\right) \\
&= \lim_{N \rightarrow \infty} \prod_{m=n}^N \left(1 - \frac{1}{m}\right) \\
&= \lim_{N \rightarrow \infty} \prod_{m=n}^N \frac{m-1}{m} \\
&= \lim_{N \rightarrow \infty} \frac{n-1}{n} \frac{n}{n+1} \cdots \frac{N-2}{N-1} \frac{N-1}{N} \\
&= \lim_{N \rightarrow \infty} \frac{n-1}{N} \\
&= 0,
\end{aligned}$$

and therefore for all $\varepsilon > 0$ that are small enough,

$$\lim_{n \rightarrow \infty} \Pr\left(\sup_{m \geq n} |X_m| < \varepsilon\right) = \lim_{n \rightarrow \infty} P_{n,\varepsilon} = 0,$$

which violates (1). In this example, almost sure convergence of X_n to zero fails as well as convergence in mean, while we still have that $X_n \rightarrow_p 0$.

Next, we show that convergence in r -th mean implies convergence in probability. The proof requires the following Lemma.

Lemma 1. (Markov's Inequality) *Let X be a random variable. For $\varepsilon > 0$ and $r > 0$,*

$$\Pr(|X| \geq \varepsilon) \leq E|X|^r / \varepsilon^r.$$

Proof. Let f_X be the PDF of X (the proof is similar for the discrete case). Let $I(\cdot)$ be an indicator function, i.e. it is equal one if the condition inside the parenthesis is satisfied, and zero otherwise. For example,

$$I(|x| \geq \varepsilon) = \begin{cases} 1, & |x| \geq \varepsilon, \\ 0, & |x| < \varepsilon. \end{cases}$$

Note that $I(|x| \geq \varepsilon) + I(|x| < \varepsilon) = 1$. Next,

$$\begin{aligned}
E|X|^r &= E[|X|^r I(|X| \geq \varepsilon)] + E[|X|^r I(|X| < \varepsilon)] \\
&\geq E[|X|^r I(|X| \geq \varepsilon)] \\
&\geq \varepsilon^r E[I(|X| \geq \varepsilon)] \\
&= \varepsilon^r \int_{-\infty}^{\infty} I(|x| \geq \varepsilon) f_X(x) dx \\
&= \varepsilon^r \left(\int_{-\infty}^{-\varepsilon} 1 \cdot f_X(x) dx + \int_{-\varepsilon}^{\varepsilon} 0 \cdot f_X(x) dx + \int_{\varepsilon}^{\infty} 1 \cdot f_X(x) dx \right) \\
&= \varepsilon^r \left(\int_{-\infty}^{-\varepsilon} f_X(x) dx + \int_{\varepsilon}^{\infty} f_X(x) dx \right) \\
&= \varepsilon^r \Pr(|X| \geq \varepsilon).
\end{aligned}$$

□

Now, suppose that X_n converges to X in r -th mean, $E |X_n - X|^r \rightarrow 0$. Then,

$$\begin{aligned} \Pr(|X_n - X| \geq \varepsilon) &\leq E |X_n - X|^r / \varepsilon^r \\ &\rightarrow 0. \end{aligned}$$

The following are some rules for manipulation of probability limits. Suppose that $X_n \rightarrow_p a$ and $Y_n \rightarrow_p b$, where a and b are some finite constants. Let c be another constant. Then,

- (i) $cX_n \rightarrow_p ca$.
- (ii) $X_n + Y_n \rightarrow_p a + b$.
- (iii) $X_n Y_n \rightarrow_p ab$.
- (iv) $X_n / Y_n \rightarrow_p a/b$, provided that $b \neq 0$.

Proof of (ii).

$$\begin{aligned} \Pr(|(X_n + Y_n) - (a + b)| \geq \varepsilon) &= \Pr(|(X_n - a) + (Y_n - b)| \geq \varepsilon) \\ &\leq \Pr(|X_n - a| + |Y_n - b| \geq \varepsilon) \\ &\leq \Pr(|X_n - a| \geq \varepsilon/2 \text{ or } |Y_n - b| \geq \varepsilon/2) \\ &\leq \Pr(|X_n - a| \geq \varepsilon/2) + \Pr(|Y_n - b| \geq \varepsilon/2) \\ &\rightarrow 0. \end{aligned}$$

□

The following result shows that if a sequence of random variables converges in probability to a constant, then their continuous functions converge in probability as well.

Theorem 2. (Slutsky's Theorem) Suppose that $X_n \rightarrow_p c$, a constant, and let $h(\cdot)$ be a continuous function at c . Then, $h(X_n) \rightarrow_p h(c)$.

Proof: By continuity of $h(\cdot)$, given $\varepsilon > 0$, there exists $\delta_\varepsilon > 0$ such that $|u - c| < \delta_\varepsilon$ implies that $|h(u) - h(c)| < \varepsilon$. Consequently, we have the following relation between the two events:

$$\{\omega : |h(X_n(\omega)) - h(c)| < \varepsilon\} \supset \{\omega : |X_n(\omega) - c| < \delta_\varepsilon\},$$

and, therefore,

$$\begin{aligned} \Pr(|h(X_n) - h(c)| < \varepsilon) &\geq \Pr(|X_n - c| < \delta_\varepsilon) \\ &\rightarrow 1. \end{aligned}$$

□

For example, suppose that $\hat{\beta}_n \rightarrow_p \beta$. Then $\hat{\beta}_n^2 \rightarrow_p \beta^2$, and $1/\hat{\beta}_n \rightarrow_p 1/\beta$, provided $\beta \neq 0$.

The random vectors/matrices converge in probability if their elements converge in probability. Alternatively, one may consider convergence in probability of norms. Consider the vector case. Let $\{X_n : n = 1, 2, \dots\}$ be a sequence of random k -vectors. We will show that $X_n \rightarrow_p 0$ element-by-element, where X is a possibly random k -vector, if and only if $\|X_n - X\| \rightarrow_p 0$, where $\|\cdot\|$ denotes the Euclidean norm. First, suppose that for all $i = 1, \dots, k$ we have that $X_{n,i} - X_i \rightarrow_p 0$. Then,

$$\begin{aligned} \|X_n - X\|^2 &= \sum_{j=1}^k (X_{n,j} - X_j)^2 \\ &\rightarrow_p 0, \end{aligned}$$

due to Slutsky's Theorem and property (ii) above. Next, suppose that $\|X_n - X\|^2 \rightarrow_p 0$. Since $\|X_n - X\|^2 = \sum_{j=1}^k (X_{n,j} - X_j)^2$, and $(X_{n,j} - X_j)^2 \geq 0$ for all n and $j = 1, \dots, k$, it must be that $X_{n,j} - X_j \rightarrow_p 0$.

The rules for manipulation of probability limits in the vector/matrix case are similar to those in the scalar case, (i) - (iv) above with corresponding definitions of multiplication and division. The Slutsky's Theorem is valid in vector/matrix case as well.

Weak Law of Large Numbers (WLLN)

The WLLN is one of the most important examples of convergence in probability.

Theorem 3. (WLLN) *Let $X_1, \dots, X_n$ be a sample of iid random variables such that $E|X_1| < \infty$. Then, $n^{-1} \sum_{i=1}^n X_i \rightarrow_p EX_1$ as $n \rightarrow \infty$.*

Note that due to iid assumption, we have that $EX_i = EX_1$ for all $i = 1, \dots, n$. We will prove the result assuming instead that $EX_1^2 < \infty$, which implies that $E|X_1| < \infty$, and $\text{Var}(X_1) < \infty$.

Theorem 4. *Let $X_1, \dots, X_n$ be a sample of iid random variables such that $\text{Var}(X_1) < \infty$. Then, $n^{-1} \sum_{i=1}^n X_i \rightarrow_p EX_1$ as $n \rightarrow \infty$.*

Proof:

$$\begin{aligned}
 \Pr \left(\left| n^{-1} \sum_{i=1}^n X_i - EX_1 \right| \geq \varepsilon \right) &= \Pr \left(\left| n^{-1} \sum_{i=1}^n (X_i - EX_1) \right| \geq \varepsilon \right) \\
 &\leq \frac{E \left| \sum_{i=1}^n (X_i - EX_1) \right|^2}{n^2 \varepsilon^2} \\
 &= \frac{\sum_{i=1}^n \sum_{j=1}^n E [X_i - EX_1] (X_j - EX_1)}{n^2 \varepsilon^2} \\
 &= \frac{\sum_{i=1}^n E [X_i - EX_1]^2}{n^2 \varepsilon^2} \\
 &= \frac{n \text{Var}(X_1)}{n^2 \varepsilon^2} \\
 &\rightarrow 0 \text{ as } n \rightarrow \infty.
 \end{aligned}$$

□

As one can see from its proof, Theorem 4 actually holds under a weaker condition than iid. For example, instead of the iid part of the assumptions of the theorem, it is sufficient to assume that $EX_i = EX_1$ for all i 's, and $\text{Cov}(X_i, X_j) = 0$ for all $i \neq j$ (uncorrelatedness). The proof then goes through without a change. In fact, those conditions can be weakened further without much change in the proof.

Convergence in distribution

Convergence in distribution is another stochastic convergence concept used to approximate the distribution of a random variable X_n in large samples. Let $\{X_n : n = 1, 2, \dots\}$ be a sequence of random variables. Let $F_n(x)$ denote the marginal CDF of X_n , i.e. $F_n(x) = \Pr(X_n \leq x)$. Let $F(x)$ be another CDF. We say that X_n *converges in distribution* if $F_n(x) \rightarrow F(x)$ for all x where $F(x)$ is continuous. In this case, we write $X_n \rightarrow_d X$, where X is *any* random variable with the distribution function $F(x)$. Note that while we say that X_n converges to X , the convergence in distribution is not convergence of random variables, but of the distribution functions.

The extension to the vector case is straightforward. Let X_n and X be two random k -vectors. We say that $X_n \rightarrow_d X$ if the joint CDF of X_n converges to that of X at all continuity points, i.e.

$$\begin{aligned}
 F_n(x_1, \dots, x_k) &= \Pr(X_{n,1} \leq x_1, \dots, X_{n,k} \leq x_k) \\
 &\rightarrow \Pr(X_1 \leq x_1, \dots, X_k \leq x_k) \\
 &= F(x_1, \dots, x_k)
 \end{aligned}$$

for all points $(x_1, \dots, x_k)$ where F is continuous. In this case, we say that the elements of $X_n, X_{n,1}, \dots, X_{n,k}$, *jointly* converge in distribution to $X_1, \dots, X_k$, the elements of X .

The rules for manipulation of convergence in distribution results are as follows.

(i) **Cramer Convergence Theorem:** Suppose that $X_n \rightarrow_d X$, and $Y_n \rightarrow_p c$. Then,

- (a) $X_n + Y_n \rightarrow_d X + c$.
- (b) $Y_n X_n \rightarrow_d cX$,
- (c) $X_n/Y_n \rightarrow_d X/c$, provided that $c \neq 0$.

Similar results hold in the vector/matrix case with proper definitions of multiplication and division.

(ii) If $X_n \rightarrow_p X$, then $X_n \rightarrow_d X$. Converse is not true with one exception:

(iii) If $X_n \rightarrow_d C$, a constant, then $X_n \rightarrow_p C$.

(iv) If $X_n - Y_n \rightarrow_p 0$, and $Y_n \rightarrow_d Y$, then $X_n \rightarrow_d Y$.

The following theorem extends convergence in distribution of random variables/vectors to convergence of their continuous functions.

Theorem 5. (Continuous Mapping Theorem (CMT)) Suppose that $X_n \rightarrow_d X$, and let $h(\cdot)$ be a function continuous on a set $\mathcal{X}$ such that $\Pr(X \in \mathcal{X}) = 1$. Then, $h(X_n) \rightarrow_d h(X)$.

Examples:

- Suppose that $X_n \rightarrow_d X$. Then $X_n^2 \rightarrow_d X^2$. For example, if $X_n \rightarrow_d N(0, 1)$, then $X_n^2 \rightarrow_d \chi_1^2$.
- Suppose that $(X_n, Y_n) \rightarrow_d (X, Y)$ (joint convergence in distribution), and set $h(x, y) = x$. Then $X_n \rightarrow_d X$. Set $h(x, y) = x^2 + y^2$. Then $X_n^2 + Y_n^2 \rightarrow_d X^2 + Y^2$. For example, if $(X_n, Y_n) \rightarrow_d N(0, I_2)$ (bivariate standard normal distribution), then $X_n^2 + Y_n^2 \rightarrow_d \chi_2^2$.

Note that contrary to convergence in probability, $X_n \rightarrow_d X$ and $Y_n \rightarrow_d Y$ does not imply that, for example, $X_n + Y_n \rightarrow_d X + Y$, unless a joint convergence result holds. This is due to the fact that the individual convergence in distribution is convergence of the *marginal* CDFs. In order to characterize the limiting distribution of $X_n + Y_n$ one has to consider the limiting behavior of the *joint* CDF of X_n and Y_n .

The Central Limit Theorem (CLT)

Various versions of the CLT are used to establish convergence in distribution of re-scaled sums of random variables.

Theorem 6. (CLT) Let $X_1, \dots, X_n$ be a sample of iid random variables such that $EX_1 = 0$ and $0 < EX_1^2 < \infty$. Then, as $n \rightarrow \infty$, $n^{-1/2} \sum_{i=1}^n X_i \rightarrow_d N(0, EX_1^2)$.

For example, the CLT can be used to approximate the distribution of the average in large samples as follows. Let $X_1, \dots, X_n$ be a sample of iid random variables with $EX_1 = \mu$ and $\text{Var}(X_1) = \sigma^2 < \infty$. Define

$$\bar{X}_n = n^{-1} \sum_{i=1}^n X_i.$$

Consider $n^{-1/2} \sum_{i=1}^n (X_i - \mu)$. We have that $(X_1 - \mu), \dots, (X_n - \mu)$ are iid with the mean $E[X_1 - \mu] = 0$, and the variance $E[X_1 - \mu]^2 = \sigma^2 < \infty$. Therefore, by the CLT,

$$\begin{aligned} n^{1/2} (\bar{X}_n - \mu) &= n^{-1/2} \sum_{i=1}^n (X_i - \mu) \\ &\rightarrow_d N(0, \sigma^2). \end{aligned}$$

In practice, we use convergence in distribution as an *approximation*. Let $\overset{a}{\sim}$ denote "approximately in large samples". Informally, one can say that $n^{1/2} (\bar{X}_n - \mu) \overset{a}{\sim} N(0, \sigma^2)$ or

$$\bar{X}_n \overset{a}{\sim} N(\mu, \sigma^2/n),$$

Note that under the normality assumption for X_i 's, the above result is obtained *exactly* for any sample size n .

The CLT can be extended to the vector case by the means of the following result.

Lemma 7. (Cramer-Wold device) Let X_n be a random k -vector. Then, $X_n \rightarrow_d X$ if and only if $\lambda^\top X_n \rightarrow_d \lambda^\top X$ for all non-zero $\lambda \in R^k$.

Corollary 8. (Multivariate CLT) Let $X_1, \dots, X_n$ be a sample of iid random k -vectors such that $EX_1 = 0$ and $EX_{1,j}^2 < \infty$ for all $j = 1, \dots, k$, and $E[X_1 X_1^\top]$ is positive definite. Then, $n^{-1/2} \sum_{i=1}^n X_i \rightarrow_d N(0, E[X_1 X_1^\top])$.

Proof: Let λ be a k -vector of constants. Consider $Y_i = \lambda^\top X_i$. We have that $Y_1, \dots, Y_n$ are iid. Further,

$$\begin{aligned} EY_1 &= \lambda^\top EX_1 \\ &= 0, \\ \text{Var}(Y_1) &= EY_1^2 \\ &= \lambda^\top E[X_1 X_1^\top] \lambda. \end{aligned}$$

The variance of Y_1 is finite provided that all the elements of the variance-covariance matrix $E[X_1 X_1^\top]$ are finite. In order to show that, note that the (r, s) -th element of $E[X_1 X_1^\top]$ is given by $E[X_{1,r} X_{1,s}]$. By the Cauchy-Schwarz inequality,

$$E|X_{1,r} X_{1,s}| \leq \sqrt{EX_{1,r}^2 EX_{1,s}^2},$$

which is finite for all $r = 1, \dots, k$, $s = 1, \dots, k$ due to the assumption that $EX_{1,j}^2 < \infty$ for all $j = 1, \dots, k$. Consequently, $EY_1^2 < \infty$, and it follows from the univariate CLT that

$$n^{-1/2} \sum_{i=1}^n Y_i \rightarrow_d N(0, \lambda^\top E[X_1 X_1^\top] \lambda).$$

Let W be any $N(0, E[X_1 X_1^\top])$ random vector. Since $\lambda^\top W \sim N(0, \lambda^\top E[X_1 X_1^\top] \lambda)$, we have that

$$\begin{aligned} \lambda^\top \left(n^{-1/2} \sum_{i=1}^n X_i \right) &= n^{-1/2} \sum_{i=1}^n Y_i \\ &\rightarrow_d \lambda^\top W. \end{aligned}$$

Therefore, by the Cramer-Wold device we have that

$$\begin{aligned} n^{-1/2} \sum_{i=1}^n X_i &\rightarrow_d W \\ &= N(0, E[X_1 X_1^\top]). \end{aligned}$$

□

Delta method

The delta method is used to derive the asymptotic distribution of the nonlinear functions of estimators. For example, in the case of iid random sample, we have that by the WLLN the average converges in probability

to the expected value of an observation: $\bar{X}_n \rightarrow_p EX_1 = \mu$. Further, it follows from Slutsky's Theorem that $h(\bar{X}_n) \rightarrow_p h(\mu)$. However, this does not allow us to approximate the distribution of $h(\bar{X}_n)$, since $h(\mu)$ is a constant (non-random). Note, that the CMT cannot be applied to general nonlinear $h(\bar{X}_n)$, since we have only a convergence in distribution result for $n^{1/2}(\bar{X}_n - \mu)$.

Theorem 9. (*Delta method*) Let $\hat{\theta}_n$ be a random k -vector, and suppose that $n^{1/2}(\hat{\theta}_n - \theta) \rightarrow_d Y$ as $n \rightarrow \infty$, where θ is a k -vector of constants, and Y is a random k -vector. Let $h : R^k \rightarrow R^m$ be a function continuously differentiable on some open neighborhood of θ . Then, $n^{1/2}(h(\hat{\theta}_n) - h(\theta)) \rightarrow_d \frac{\partial h(\theta)}{\partial \theta^\top} Y$.

Proof: First, note that $n^{1/2}(\hat{\theta}_n - \theta) \rightarrow_d Y$ implies that $\hat{\theta}_n - \theta \rightarrow_p 0$ or $\hat{\theta}_n \rightarrow_p \theta$. Indeed, define $\tau_n = n^{-1/2}$. We have that $\tau_n \rightarrow 0$, and, consequently, $\tau_n \rightarrow_p 0$. By the Cramer Convergence Theorem (b),

$$\begin{aligned} (\hat{\theta}_n - \theta) &= \tau_n n^{1/2} (\hat{\theta}_n - \theta) \\ &\rightarrow_d (\text{plim } \tau_n) Y \\ &= 0. \end{aligned}$$

Therefore, by property (iii) of convergence in distribution, $\hat{\theta}_n \rightarrow_p \theta$.

Apply the mean value theorem to the function $h(\hat{\theta}_n)$ element-by-element (see the Appendix) to obtain

$$h(\hat{\theta}_n) = h(\theta) + \frac{\partial h(\theta_n^*)}{\partial \theta^\top} (\hat{\theta}_n - \theta), \quad (3)$$

where θ_n^* is a random variable that lies between $\hat{\theta}_n$ and θ (element-by-element), i.e. $\|\theta_n^* - \theta\| \leq \|\hat{\theta}_n - \theta\|$. Since $\hat{\theta}_n \rightarrow_p \theta$, it has to be that $\theta_n^* \rightarrow_p \theta$ as well:

$$\begin{aligned} \Pr(\|\theta_n^* - \theta\| \geq \varepsilon) &\leq \Pr(\|\hat{\theta}_n - \theta\| \geq \varepsilon) \\ &\rightarrow 0. \end{aligned}$$

Furthermore, by Slutsky's Theorem,

$$\frac{\partial h(\theta_n^*)}{\partial \theta^\top} \rightarrow_p \frac{\partial h(\theta)}{\partial \theta^\top}. \quad (4)$$

Next, re-write (3) as follows:

$$n^{1/2}(h(\hat{\theta}_n) - h(\theta)) = \frac{\partial h(\theta_n^*)}{\partial \theta^\top} n^{1/2}(\hat{\theta}_n - \theta).$$

Then, it follows from the result in (4), assumption $n^{1/2}(\hat{\theta}_n - \theta) \rightarrow_d Y$ and Cramer Convergence Theorem (b) that

$$n^{1/2}(h(\hat{\theta}_n) - h(\theta)) \rightarrow_d \frac{\partial h(\theta)}{\partial \theta^\top} Y.$$

□

Consider again the example of the average of iid random variables with finite variance. We have that $n^{1/2}(\bar{X}_n - \mu) \rightarrow_d N(0, \sigma^2)$. Suppose that $\mu \neq 0$. Then, by the delta method,

$$\begin{aligned} n^{1/2}\left(\frac{1}{\bar{X}_n} - \frac{1}{\mu}\right) &\rightarrow_d -\frac{1}{\mu^2}N(0, \sigma^2) \\ &= N\left(0, \frac{\sigma^2}{\mu^4}\right). \end{aligned}$$

Appendix A: Mean value theorem

Theorem 10. (*One-Dimensional Mean-Value Theorem*) Let $f : [a, b] \rightarrow R$ be continuous on $[a, b]$ and differentiable on (a, b) . Then there is $c \in (a, b)$ such that

$$f(b) - f(a) = \frac{df(c)}{dx} (b - a).$$

Now, suppose we have $h : \Theta \rightarrow R$, where $\Theta \subset R^k$. Suppose further that h is continuously differentiable on some open neighborhood of θ_0 , say N_0 , and let u be such that $\theta_0 + tu \in N_0$ for all $t \in [0, 1]$. Define $f(t) = h(\theta_0 + tu)$. The function f is continuous and differentiable on $[0, 1]$ interval, and by the one-dimensional mean-value theorem,

$$\begin{aligned} h(\theta_0 + u) - h(\theta_0) &= f(1) - f(0) \\ &= \frac{df(t^*)}{dt} \text{ for some } t^* \in (0, 1) \\ &= \frac{\partial h(\theta_0 + t^*u)}{\partial \theta^\top} u \\ &= \frac{\partial h(\theta^*)}{\partial \theta^\top} u, \end{aligned}$$

where

$$\theta^* = \theta_0 + t^*u,$$

and

$$\begin{aligned} \|\theta^* - \theta_0\| &= t^* \|u\| \\ &< \|u\|. \end{aligned}$$

(The argument follows closely that of Theorem 10 on page 106 of Magnus and Neudecker (2007): *Matrix Differential Calculus with Applications in Statistics and Econometrics*, 3rd Edition.) We have established a mean-value theorem for real-valued functions of several variables:

Theorem 11. Let $h : \Theta \rightarrow R$, where $\Theta \subset R^k$, be continuously differentiable on some open neighborhood N_θ of θ . If $\hat{\theta} \in N_\theta$, then there is $\theta^* \in N_\theta$ such that

$$h(\hat{\theta}) - h(\theta) = \frac{\partial h(\theta^*)}{\partial \theta^\top} (\hat{\theta} - \theta),$$

where $\|\theta^* - \theta\| \leq \|\hat{\theta} - \theta\|$.

If $h(\theta) = (h_1(\theta), \dots, h_m(\theta))^\top$ is a vector valued function with $h_j : \Theta \rightarrow R$ for all $j = 1, \dots, m$, the above theorem can be applied element-by-element:

$$\begin{aligned} h(\hat{\theta}) - h(\theta) &= \begin{pmatrix} h_1(\hat{\theta}) - h_1(\theta) \\ \vdots \\ h_m(\hat{\theta}) - h_m(\theta) \end{pmatrix} \\ &= \begin{pmatrix} \frac{\partial h(\theta^{*,1})}{\partial \theta^\top} (\hat{\theta} - \theta) \\ \vdots \\ \frac{\partial h(\theta^{*,m})}{\partial \theta^\top} (\hat{\theta} - \theta) \end{pmatrix} \\ &= \begin{pmatrix} \frac{\partial h(\theta^{*,1})}{\partial \theta^\top} \\ \vdots \\ \frac{\partial h(\theta^{*,m})}{\partial \theta^\top} \end{pmatrix} (\hat{\theta} - \theta), \end{aligned}$$

where

$$\|\theta^{*,j} - \theta\| \leq \|\hat{\theta} - \theta\|, \quad (5)$$

for all $j = 1, \dots, m$. To simplify the notation, we can write

$$\begin{pmatrix} \frac{\partial h(\theta^{*,1})}{\partial \theta^\top} \\ \vdots \\ \frac{\partial h(\theta^{*,m})}{\partial \theta^\top} \end{pmatrix} = \frac{\partial h(\theta^*)}{\partial \theta^\top},$$

indicating that θ^* may be different across the rows of the matrix $\partial h(\theta^*)/\partial \theta^\top$, and that, in each row, θ^* satisfies (5).

Appendix B: Proof of the CLT

The material discussed here is adapted from Hogg, McKean, and Craig (2005): *Introduction to Mathematical Statistics*. Let $\bar{X}_n = n^{-1} \sum_{i=1}^n X_i$, where X_i 's are iid with mean μ and variance σ^2 . Recall from Lecture 1 that the MGF of a $N(0, \sigma^2)$ distribution is given by $\exp(t^2 \sigma^2 / 2)$. In view of Lemma 2 in Lecture 1, it suffices to show that the MGF of $n^{1/2}(\bar{X}_n - \mu)$ converges to $\exp(t^2 \sigma^2 / 2)$.²

Let $m(t)$ denote the MGF of $X_1 - \mu$:

$$m(t) = E \exp(t(X_1 - \mu)).$$

Recall from Lecture 1 that,

$$\begin{aligned} m(0) &= 1, \\ m^{(1)}(0) &= E[X_1 - \mu] = 0, \\ m^{(2)}(0) &= E[(X_1 - \mu)^2] = \sigma^2, \end{aligned}$$

where

$$m^{(s)}(0) = \left. \frac{d^s m(t)}{dt^s} \right|_{t=0}.$$

We have the following expansion of $m(t)$:

$$\begin{aligned} m(t) &= m(0) + m^{(1)}(0)t + \frac{m^{(2)}(s)t^2}{2} \\ &= 1 + \frac{m^{(2)}(s)t^2}{2}. \end{aligned} \quad (6)$$

where s is a mean value such that lies between 0 and t .

²If the MGF does not exist, one can replace it with the *characteristic function* of $X_1 - \mu$, which is defined as $\varphi(t) = E \exp(it(X_1 - \mu))$, where $i = \sqrt{-1}$. Note that the characteristic function always exists, and the proof with the characteristic function is essentially the same as the proof that uses the MGF.

Let $M_n(t)$ denote the MGF of $n^{1/2}(\bar{X}_n - \mu)$. We have:

$$\begin{aligned}
M_n(t) &= E \exp \left(t \frac{1}{n^{1/2}} \sum_{i=1}^n (X_i - \mu) \right) \\
&= E \prod_{i=1}^n \exp \left(\frac{t}{n^{1/2}} (X_i - \mu) \right) \\
&= \prod_{i=1}^n E \exp \left(\frac{t}{n^{1/2}} (X_i - \mu) \right) \quad (\text{by independence}) \\
&= \left(E \exp \left(\frac{t}{n^{1/2}} (X_1 - \mu) \right) \right)^n \quad (\text{because of ident. distr.}) \\
&= \left(m \left(\frac{t}{n^{1/2}} \right) \right)^n \quad (\text{by definition of } m(t)) \\
&= \left(1 + \frac{m^{(2)}(s) \left(\frac{t}{n^{1/2}} \right)^2}{2} \right)^n \quad (\text{by (6)}) \\
&= \left(1 + \frac{m^{(2)}(s) t^2}{2n} \right)^n \\
&= \left(1 + \frac{a_n}{n} \right)^n,
\end{aligned}$$

where s lies between 0 and $t/n^{1/2}$ and therefore converges to zero as $n \rightarrow \infty$, and

$$\begin{aligned}
a_n &= \frac{m^{(2)}(s) t^2}{2} \\
&\rightarrow \frac{m^{(2)}(0) t^2}{2} \quad (\text{as } n \rightarrow \infty) \\
&= \frac{\sigma^2 t^2}{2}.
\end{aligned}$$

We will show next that

$$\log M_n(t) = n \log \left(1 + \frac{a_n}{n} \right) \rightarrow \lim_{n \rightarrow \infty} a_n = \frac{\sigma^2 t^2}{2}. \quad (7)$$

Note that the result in (7) implies that

$$M_n(t) = \exp(\log M_n(t)) \rightarrow \exp \left(\lim_{n \rightarrow \infty} \log M_n(t) \right) = \exp \left(\frac{\sigma^2 t^2}{2} \right).$$

To show (7), write

$$\begin{aligned}
\lim_{n \rightarrow \infty} n \log \left(1 + \frac{a_n}{n} \right) &= \lim_{n \rightarrow \infty} \frac{\log(1 + a_n/n)}{1/n} \\
&= \lim_{n \rightarrow \infty} a_n \frac{\log(1 + a_n/n)}{a_n/n} \\
&= \lim_{n \rightarrow \infty} a_n \lim_{\delta \rightarrow 0} \frac{\log(1 + \delta)}{\delta} \quad (\text{by change of variable } \delta = a_n/n) \\
&= \frac{\sigma^2 t^2}{2} \lim_{\delta \rightarrow 0} \frac{\log(1 + \delta)}{\delta}.
\end{aligned}$$

Lastly, by l'Hôpital's rule,

$$\lim_{\delta \rightarrow 0} \frac{\log(1 + \delta)}{\delta} = \lim_{\delta \rightarrow 0} \frac{1/(1 + \delta)}{1} = 1.$$